

Assignment 4

Loic Roldan Waals

April 16, 2018

Question 2

First the data is imported

```
#determine where files are read and written
setwd("C:/Users/Loic RW/Google Drive/Big Data and Business Analytics/Assignments/Assignment 4")

#read all the data
library(readr)
tweets <- read_delim("C:/Users/Loic RW/Google Drive/Big Data and Business Analytics/Assignments/Assignment 4",
                    ";", escape_double = FALSE, col_types = cols(favorited = col_skip(),
                                                                favoriteCount = col_skip(), isRetweet = col_skip(),
                                                                latitude = col_skip(), longitude = col_skip(),
                                                                replyToSID = col_skip(), replyToSN = col_skip(),
                                                                replyToUID = col_skip(), retweeted = col_skip(),
                                                                statusSource = col_skip(), truncated = col_skip()),
                    trim_ws = TRUE)
```

```
## Warning in rbind(names(probs), probs_f): number of columns of result is not
## a multiple of vector length (arg 1)
```

```
## Warning: 4 parsing failures.
## row # A tibble: 4 x 5 col      row col   expected   actual   file
```

```
tweets$retweetCount <- as.numeric(tweets$retweetCount)
```

```
## Warning: NAs introduced by coercion
```

```
#delete 4 NA's for "retweetcount"
tweets <- tweets[complete.cases(tweets),]
```


Next we make a histogram of the number of words of the tweets.

```
library(ngram)

for (i in 1:nrow(tweets)) {
  tweets$wordCount[i] <- as.numeric(wordcount(tweets$text[i], sep = " ", count.function = sum))
}
```


```
## Warning: Unknown or uninitialised column: 'wordCount'.
```

```
hist(tweets$wordCount)
```



Next we plot number of words and number of retweets

```
plot(tweets$wordCount, tweets$retweetCount, xlab = "word count", ylab = "retweet count")
```

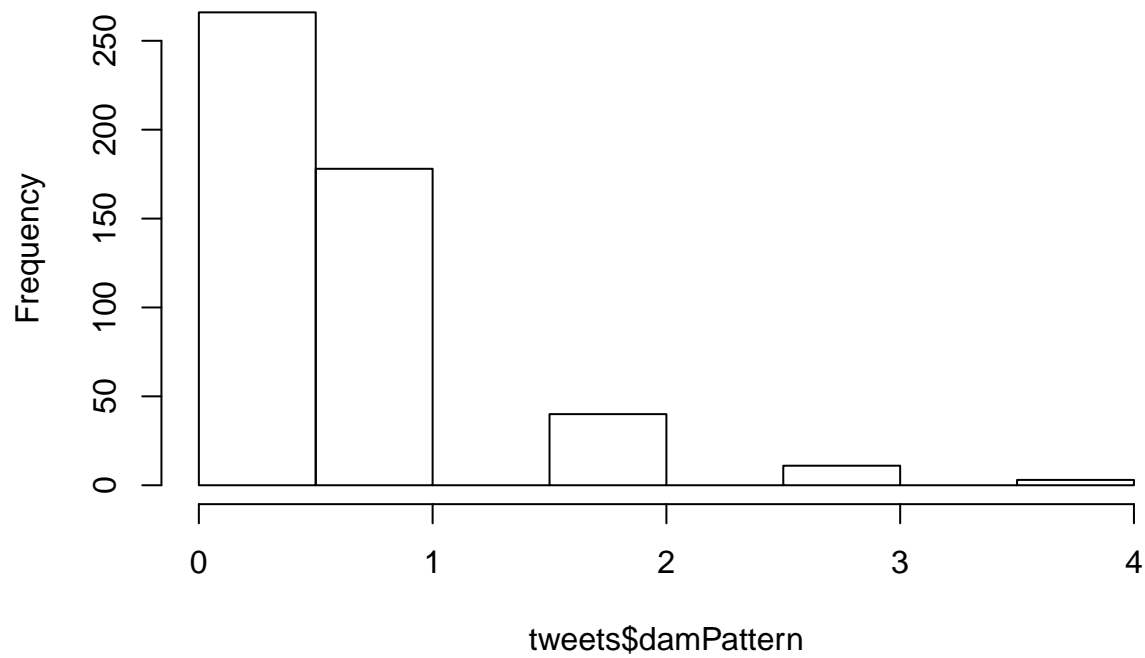


Next we calculate the frequency of the pattern “dam” per tweet. As we can see there are only two instances where the pattern “dam” is used. As this would not make for an exciting graph, we used the string “ing” instead for plotting it against time

```
library(stringr)


tweets$damPattern <- str_count(tweets$text, "ing")
hist(tweets$damPattern)
```

Histogram of tweets\$damPattern



```
tweets$realTime <- as.POSIXct(tweets$created, format = "%y-%m-%d %H:%M:%S")
tweets$minute <- as.POSIXlt(tweets$realTime)$min

tempdf <- tweets[,c("damPattern", "minute")]
tempdf <- aggregate(cbind(damPattern)~minute, data = tempdf, FUN = sum)
plot(tempdf$minute,tempdf$damPattern)
```



Question 3

First we import the data

```
library(readr)
trainTitanic <- read_csv("C:/Users/Loic RW/Google Drive/Big Data and Business Analytics/Workshops/Session 3/trainTitanic.csv",
  col_types = cols(Name = col_skip(), PassengerId = col_skip()))

#fix variables
trainTitanic$Survived <- as.ordered(trainTitanic$Survived)
trainTitanic$Pclass <- as.factor(trainTitanic$Pclass)
trainTitanic$Sex <- as.factor(trainTitanic$Sex)
trainTitanic$Embarked <- as.factor(trainTitanic$Embarked)

trainTitanic$Age[is.na(trainTitanic$Age)] <- mean(trainTitanic$Age, na.rm = TRUE)
trainTitanic$Cabin <- NULL
trainTitanic$Ticket <- NULL

trainTitanic <- trainTitanic[complete.cases(trainTitanic),]

trainTitanic$Age <- scale(trainTitanic$Age)
trainTitanic$Fare <- scale(trainTitanic$Fare)
```

Then we perform the analysis

```

library(C50)
library(e1071)

mdl <- Survived ~ .

#Xval
#cross validation
nFolds = 10
myFolds <- cut(seq(1,nrow(trainTitanic)),
               breaks = nFolds,
               labels = FALSE)

#initialise accuracy variables
accNB <- rep(NA, nFolds)
accSVM <- rep(NA, nFolds)

specNB <- rep(NA, nFolds)
specSVM <- rep(NA, nFolds)

sensNB <- rep(NA, nFolds)
sensSVM <- rep(NA, nFolds)

AUCNB <- rep(NA, nFolds)
AUCSVM <- rep(NA, nFolds)

library(caret)

## Loading required package: lattice

## Loading required package: ggplot2

library(pROC)

## Type 'citation("pROC")' for a citation.

##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##      cov, smooth, var

for (i in 1:nFolds) {
  cat("Analysis of fold", i, "\n")

  #define training and test set (print commands are to help troubleshoot
#and determine where the programme aborted during errors)
  testIndex <- which(myFolds == i, arr.ind = TRUE)
  crossTest <- trainTitanic[testIndex, ]
  crossTrain <- trainTitanic[-testIndex, ]
  print("data allocated")

```

```

#train the models
rsltNB <- naiveBayes(mdl, data = crossTrain)
print("Naive Bayes trained")
rsltSVM <- svm(mdl, data = crossTrain)
print("svm trained")

#predict values
pdNB <- as.ordered(predict(rsltNB, crossTest, type = "class"))
print("nb predicted")
pdSVM = as.ordered(predict(rsltSVM, crossTest))
print("svm predicted")
pdNB
pdSVM

#measure accuracy
accNB[i] = mean(pdNB == crossTest$Survived)
print("nb accuracy saved")
accSVM[i] = mean(pdSVM == crossTest$Survived)
print("SVM accuracy saved")

confMatrix <- table(pdNB, crossTest$Survived)
specNB[i] = specificity(confMatrix)
sensNB[i] = sensitivity(confMatrix)
print("nb other measures done")

confMatrix <- table(pdSVM, crossTest$Survived)
specSVM[i] = specificity(confMatrix)
sensSVM[i] = sensitivity(confMatrix)
print("svm other measures done")

#AUC
crossTrain$Survived <- as.numeric(crossTrain$Survived)
crossTest$Survived <- as.numeric(crossTest$Survived)

rsltNB <- naiveBayes(mdl, data = crossTrain)
print("Naive Bayes trained")
rsltSVM <- svm(mdl, data = crossTrain)
print("svm trained")

pdNB <- predict(rsltNB, crossTest, type = "raw")
pdNB <- pdNB[,2]
print("nb predicted")
pdSVM = predict(rsltSVM, crossTest)
print("svm predicted")

AUCNB[i] = as.numeric(auc(crossTest$Survived, pdNB))
AUCSVM[i] = as.numeric(auc(crossTest$Survived, pdSVM))
}

## Analysis of fold 1
## [1] "data allocated"

```

```

## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## [1] "nb accuracy saved"
## [1] "SVM accuracy saved"
## [1] "nb other measures done"
## [1] "svm other measures done"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## Analysis of fold 2
## [1] "data allocated"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## [1] "nb accuracy saved"
## [1] "SVM accuracy saved"
## [1] "nb other measures done"
## [1] "svm other measures done"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## Analysis of fold 3
## [1] "data allocated"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## [1] "nb accuracy saved"
## [1] "SVM accuracy saved"
## [1] "nb other measures done"
## [1] "svm other measures done"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## Analysis of fold 4
## [1] "data allocated"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## [1] "nb accuracy saved"
## [1] "SVM accuracy saved"
## [1] "nb other measures done"
## [1] "svm other measures done"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"

```



```
## Analysis of fold 5
## [1] "data allocated"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## [1] "nb accuracy saved"
## [1] "SVM accuracy saved"
## [1] "nb other measures done"
## [1] "svm other measures done"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## Analysis of fold 6
## [1] "data allocated"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## [1] "nb accuracy saved"
## [1] "SVM accuracy saved"
## [1] "nb other measures done"
## [1] "svm other measures done"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## Analysis of fold 7
## [1] "data allocated"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## [1] "nb accuracy saved"
## [1] "SVM accuracy saved"
## [1] "nb other measures done"
## [1] "svm other measures done"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## Analysis of fold 8
## [1] "data allocated"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## [1] "nb accuracy saved"
## [1] "SVM accuracy saved"
## [1] "nb other measures done"
## [1] "svm other measures done"
## [1] "Naive Bayes trained"
## [1] "svm trained"
```

```

## [1] "nb predicted"
## [1] "svm predicted"
## Analysis of fold 9
## [1] "data allocated"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## [1] "nb accuracy saved"
## [1] "SVM accuracy saved"
## [1] "nb other measures done"
## [1] "svm other measures done"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## Analysis of fold 10
## [1] "data allocated"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"
## [1] "nb accuracy saved"
## [1] "SVM accuracy saved"
## [1] "nb other measures done"
## [1] "svm other measures done"
## [1] "Naive Bayes trained"
## [1] "svm trained"
## [1] "nb predicted"
## [1] "svm predicted"

```

```
#determine the average accuracy over the 10 folds for each model
```

```
avgAccNB = mean(accNB)
```

```
avgAccSVM = mean(accSVM)
```

```
avgSpecNB = mean(specNB)
```

```
avgSpecSVM = mean(specSVM)
```

```
avgSensNB = mean(sensNB)
```

```
avgSensSVM = mean(sensSVM)
```

```
avgAUCNB = mean(AUCNB)
```

```
avgAUCSVM = mean(AUCSVM)
```

```
#print all the results
```

```
avgAccNB
```

```
## [1] 0.7761747
```

```
avgAccSVM
```

```
## [1] 0.8245148
```

```
avgSpecNB
```

```
## [1] 0.5849951
```

```
avgSpecSVM
```

```
## [1] 0.7150887
```

```
avgSensNB
```

```
## [1] 0.8963486
```

```
avgSensSVM
```

```
## [1] 0.8898537
```

```
avgAUCNB
```

```
## [1] 0.823757
```

```
avgAUCSVM
```

```
## [1] 0.8396036
```

Now we will examine the ROC

```
#Xval
```

```
library(caret)
```

```
library(ROSE)
```

```
## Loaded ROSE 0.0-3
```

```
library(caTools)
```

```
set.seed(12345)
```

```
trainTitanic$Survived <- as.numeric(trainTitanic$Survived)
```

```
sample <- sample.split(trainTitanic$Survived, SplitRatio = 0.8)
```

```
train = subset(trainTitanic, sample == TRUE)
```

```
test = subset(trainTitanic, sample == FALSE)
```

```
nrow(test)
```

```
## [1] 178
```

```

#train the models
rsltNB <- naiveBayes(mdl, data = train)
print("Naive Bayes trained")

## [1] "Naive Bayes trained"

rsltSVM <- svm(mdl, data = train)
print("svm trained")

## [1] "svm trained"

#predict values
pdNB <- predict(rsltNB, test, type = "raw")
pdNB <- pdNB[,2]
print("nb predicted")


## [1] "nb predicted"

pdSVM = predict(rsltSVM, test)
print("svm predicted")

## [1] "svm predicted"


roc.curve(test$Survived, pdNB)

```



```
## Area under the curve (AUC): 0.856
```

```
roc.curve(test$Survived, pdSVM)
```



```
## Area under the curve (AUC): 0.837
```